

Table of Contents

Preface	v
1 Introduction	1
1.1 Chapter Objectives	1
1.2 Why TinyML Feels Intimidating (But Shouldn't)	1
1.3 What This Book Is (and Is Not)	2
1.4 TinyML	3
1.5 Edge AI vs. TinyML	7
1.6 Constraints Are the Best Part	9
1.7 What You Will Build in This Book	9
2 Your First TinyML App	11
2.1 Chapter Objectives	11
2.2 What You Are Building: The "Shake Detector"	11
2.3 Hardware Verification: Checking the Senses	13
2.4 Running Your First TinyML Model	15
2.5 Common Issues and Fixes	21
3 Sensors, Data, and Labels	23
3.1 Chapter Objectives	23
3.2 What is a Sensor, Really?	23
3.3 What is a Data Window?	26
3.4 The Core TinyML Engineering Balancing Act	27
3.5 Labels: The Guardrails of Machine Learning	29
3.6 Dataset Balance: Avoiding Systematic Bias	30
3.7 Variability is Your Friend	30
3.8 Watching the Slicer in Action	31
3.9 Common Issues and Fixes	32
4 Edge Impulse — Basic Setup	34
4.1 Chapter Objectives	34
4.2 Why TinyML Requires a Specialized Workbench	34
4.3 Edge Impulse: The Engineering Environment	35
4.4 Setting Up Your Project Workspace	39

4.5 Connecting Your Hardware 44

4.6 Common Issues and Fixes 51

5 TinyML Data Collection and Preparation 55

5.1 Chapter Objectives 55

5.2 The Three Phases of Edge Impulse 55

5.3 Project Blueprint: The "Shake" Classifier 56

5.4 Step-by-Step Data Collection 57

5.5 Hardware and Platform Boundaries 63

5.6 Common Issues and Fixes 63

6 Training TinyML Models 65

6.1 Chapter Objectives 65

6.2 Demystifying the Training Phase 65

6.3 What Training Really Means 66

6.4 Architectural Boundaries and Constraints 68

6.5 Configuring and Running Your Impulse Pipeline 69

6.6 Common Issues and Fixes 78

7 From Model to Microcontroller (Arduino) 81

7.1 Chapter Objectives 81

7.2 What Deployment Means in TinyML 81

7.3 Deploy the Model 84

7.4 Including the Model in Your Program 88

7.5 Matching Runtime Data to Training Data 88

7.6 Running Inference 89

7.7 Memory and Performance Reality Check 90

7.8 Deployment Is an Iteration Step 91

7.9 Common Issues and Fixes 91

8 From Model to Microcontroller (C++) 94

8.1 Chapter Objectives 94

8.2 Arduino Vs C++ Inference Library 95

8.3 Building the C++ Library From Edge Impulse 95

8.4 Using the C++ Library Directly (Practical) 99

8.5 Common Issues and Fixes 101

9 Making Predictions Reliable 104

9.1 Chapter Objectives 104

9.2 Why "Working" Isn't Always Good Enough 104

9.3 Using Confidence Thresholds 105

9.4 Smoothing Predictions Over Time 106

9.5 Timing Matters (Don't React Too Often) 108

9.6 Separating Recognition From Action 108

9.7 A Complete Reliable Pipeline 108

9.8 Common Issues and Fixes 109

10	Building Efficient and Stable Models	111
10.1	Chapter Objectives	111
10.2	From Working Models to Real Systems	111
10.3	Creating Project and Connecting Device	113
10.4	Data Acquisition	113
10.5	Create Impulse And Generate Features	115
10.6	Train And Deploy TinyML Model	116
10.7	Understanding Model Size	118
10.8	Balancing Accuracy and Memory	122
10.9	Measuring Inference Performance	124
10.10	Preventing Instability	126
10.11	Common Issues and Fixes	127
11	TinyML Model With Sensor Fusion	130
11.1	Chapter Objectives	130
11.2	Why Sensor Fusion Matters	130
11.3	Case Study: Gesture Recognition	131
11.4	Dataset Collection	135
11.5	Data Cleaning (Critical Step)	135
11.6	Generate Features	141
11.7	Develop and Train a TinyML Model	143
11.8	Deployment	144
11.9	Common Issues and Fixes	146
12	From TinyML to Edge AI	149
12.1	Chapter Objectives	149
12.2	From TinyML to Edge AI	149
12.3	Why Smartphones Matter	150
12.4	Running Models on Smartphones	152
12.5	Example: Motion Recognition on a Smartphone	153
12.6	Example: Sound Recognition on a Smartphone	162
A	Microcontrollers for TinyML	172